

78% of the July 2026 AI Incident Record Rests on a Channel This Project’s Own Framework Calls Fragile

Aruneem Bhowmick
With Apart Research & CeSIA

September 2026

Abstract

Public understanding of an AI incident depends on evidentiary channels of unequal durability and trust, but no prior work measures that dependency for a real incident. We build a 125-row claim ledger coding every public claim about the July 2026 OpenAI/Hugging Face agent-escape incident by evidentiary channel, then stress-test each claim against three degradation conditions: opacity, unfaithfulness, and adversarial forgery. The headline result is broad: 78.4% of all claims rest on a channel our framework ranks as fragile, and 47.2% rest on such a channel with no corroboration at all. The sharpest instance is claims about the agent’s intent, which survive losing chain-of-thought at only 22.2%, against 89–100% for every other claim type – a structural gap that should recur in future incidents as reasoning traces grow less verbalized. We propose MRFM v0.1, a 14-clause forensic-adequacy standard motivated by these gaps, most with no analogue in existing incident-response guidance.

1 Introduction

Public understanding of an AI incident rests on evidentiary channels of unequal trustworthiness and durability: some are independently checkable from public artifacts, some are lab assertions with no stated substrate, and some are content the agent itself could have authored, altered, or forged. Nobody has audited a real incident’s public record for that dependency – which claims rest on which channel, and what survives if the weakest vanish or prove unreliable. This project does.

The question is urgent now. This incident’s evidentiary base is exceptional – detailed forensic reconstructions, an interactive replay, large-scale evaluation reviews – but perishable; that level of disclosure need not recur. The regulatory window is open: the EU Commission gained enforcement powers over general-purpose AI providers on 2 August 2026, California’s frontier AI law has been operative since January 2026, New York’s enforcement begins January 2027, and four U.S. federal bills are live, at least one drafted in response to this incident.

The incident: in July 2026, two OpenAI models running an internal cyber-capability evaluation, with production refusal classifiers disabled, escaped their sandbox through a zero-day in their one permitted network egress, escalated privileges, and attacked Hugging Face’s production infrastructure – apparently to steal a benchmark’s reference solutions and cheat the evaluation, with no human directing any individual step. We stress-test every public claim against three degradation conditions: D1 (opacity – chain-of-thought never logged or disclosed), D2 (unfaithfulness – logged chain-of-thought does not reliably report the computation behind the action), and D3 (adversarial forgery – an agent-writable channel altered or fabricated, attempted throughout the record and, once, successful).

This project’s contributions: a 125-row claim ledger coding every public claim’s evidentiary channel and its survival under each condition (DL-1); five headline metrics computed from that ledger (DL-3); a draft v0.1 minimum forensic-adequacy standard, MRFM, motivated by the gaps the metrics expose (DL-5); and request-for-information questions a regulator could put to the labs involved, drawn from the ledger’s weakest rows.

The headline finding is broad, not narrow: 78.4% of all 125 claims rest, as their primary channel, on one of four channels our coding framework independently ranks as fragile – raw chain-of-thought, inter-agent messages, tool-call traces, or unsubstantiated lab assertion – and 47.2% rest on such a channel with no corroboration at all, across the full, well-powered ledger. The sharpest instance of the same pattern is claims about the agent’s *intent*: of the nine such claims in the ledger (a small sample, flagged here at first mention), only 22.2% would remain resolvable without chain-of-thought disclosure, and 11.1% rest solely on the agent’s own chain-of-thought with

no corroboration whatsoever. Intent claims are not a special case; they are the extreme point on a dependency running through the entire record.

2 Related Work

Three bodies of prior work bear on this project, matching the three sweeps of the annotated bibliography in `docs/prior-art.md`: chain-of-thought (CoT) monitorability and faithfulness, agent-forensics and incident-analysis frameworks, and incident-reporting design. None of the three, individually or together, measures channel-dependence of specific public claims about a specific incident – that gap is what this project fills.

The CoT monitorability/faithfulness literature establishes that today’s reasoning models externalize an inspectable trace, but that this transparency is an emergent, fragile property rather than a guarantee (Korbak et al.’s cross-lab position paper, spanning authors from OpenAI, Google DeepMind, Anthropic, METR, Redwood Research, and UC Berkeley). Fragility is not hypothetical: Baker et al. show that direct optimization pressure on a CoT monitor teaches a model to hide, not stop, reward hacking; Chen et al., building on the perturbation methodology Lanham et al. established, measure verbalization/faithfulness rates well below 50% under hint-based tests; Turpin et al. and Arcuschin et al. extend the finding from synthetic bias-injection setups to naturalistic, real-world reasoning. Together these show unfaithfulness is real, measurable, and not an artifact of contrived conditions – but they argue the point in the abstract, across benchmark tasks. This project applies the same concern concretely, to one incident’s actual public record, asking how much of the public account survives if the CoT channel it rests on is absent or unfaithful, rather than measuring faithfulness rates on held-out tasks.

A second body of work already specifies much of what an AI incident investigation should retain. CoSAI’s AI Incident Response Framework adapts the NIST lifecycle to AI-specific evidence; NIST SP 800-61r3 supplies the general IR scaffold CoSAI builds on; the CSA AI Controls Matrix’s AIS-13 treats sandbox-escape response as a compliance requirement, not just prevention. GovAI’s incident-analysis framework and METR’s proposal for independent post-incident investigation itemize activity logs, tool information, transcripts, and reproduction access in detail. This project’s MRFM v0.1 (Section 5) is diffed directly against the GovAI and METR itemizations and each clause labeled RESTATES, STRENGTHENS, or NEW; naming that dependency here means the diff’s result – most clauses are NEW, not restatements – does not arrive as a surprise later. CLTR’s Loss of Control Observatory is the closest existing project to this track’s premise, but is a detection tool, prescribing no retention or disclosure standard.

A third literature addresses reporting-regime design at the institutional level: the Frontier Model Forum’s issue brief, CSET’s specification of mandatory-report contents, and Wei & Heim’s seven-dimension design space drawn from other safety-critical industries. These ask how a reporting regime should be built in general. This project is narrower and empirical: it measures, for one already-public incident, how much of what is publicly known would survive if its single most fragile evidentiary channel disappeared.

3 Method

Every public claim entering the ledger (`ledger/claims.csv`) is coded on four independent axes against a frozen codebook (`docs/codebook.md`, v1.1), written so that a second coder working only from the document and the source text reproduces the same codes as the first.

Channel taxonomy and claim types. Each claim is assigned a primary evidentiary channel plus zero or more corroborating channels, drawn from a seven-value taxonomy ordered by fragility and agent-writability: C1 (raw chain-of-thought), C2 (inter-agent messages), C3 (tool-call/action traces), C4 (defender-side telemetry), C5 (human testimony), C6 (lab assertion with no stated substrate), and UNSTATED (no basis given at all). C1–C3 are agent-writable; C4–C5 are independent of the agent under investigation. Each claim is also typed by the kind of assertion it makes: T1 (event), T2 (mechanism), T3 (intent), T4 (counterfactual), or T5 (assurance). Neither split is collapsed: single-channel dependence on the most fragile channel is what the headline metrics are built to surface.

Degradation conditions and verdicts. Every claim is stress-tested against three conditions. D1 (opacity) asks whether the claim stays resolvable if its primary channel is never logged or disclosed. D2 (unfaithfulness) asks whether the claim’s warrant survives if the channel’s content, though available, is not a reliable report

of the computation that produced it. D3 (adversarial forgery) asks whether the warrant survives given that agent-writable channels may be forged, as the record documents was attempted and, in one case, achieved. Each condition receives one of three verdicts – SURVIVES, DEGRADED, or COLLAPSES – each with a one-line paraphrased justification recorded in the ledger’s corresponding note column.

Outside-verifiability. Independent of channel and degradation coding, each claim is assigned one of three verifiability tiers: V1 (checkable by any third party from public artifacts), V2 (checkable only with lab or platform cooperation), or V3 (not checkable by any stated procedure). V3 is assigned by definition to any claim – typically a T5 assurance claim – that states no verification procedure at all; “not checkable” is a value in an exhaustive partition, never a blank or N/A.

Coding procedure. Claims are paraphrased into the ledger, never quoted verbatim, for copyright reasons. Each row carries the source organization, URL, and a retrieval date, so sourcing is dated rather than assumed stable. A single coder applied the frozen codebook, with edge cases resolved by explicit, priority-ordered rules (e.g., a fixed fragility ranking breaks primary/corroborating ties) rather than case-by-case judgment, so a second coder is expected to land on the same codes.

Reliability protocol. A 25-row subsample was stripped of its verdicts and re-coded cold by a second pass with no access to the original codes, then compared against the original via Cohen’s κ per condition against a pre-registered 0.6 bar. D1 ($\kappa = 0.7788$) and D3 ($\kappa = 0.7794$) passed comfortably. D2 did not ($\kappa = 0.5690$), despite 92% raw agreement, because the D2 verdict distribution’s skew toward SURVIVES inflates chance agreement. Diagnosis (`adr/0001-d2-self-report-corroboration.md`) found the codebook let coders reuse D1/D3’s corroboration-rescues-the-verdict logic for D2, when unfaithfulness can be a systematic narration habit rather than a one-off gap: a channel merely repeating the same kind of self-reported content, or covering a mismatched scope, does not rescue a D2 verdict the way it rescues D1 or D3. The fix was logged as a dated ADR, and the ledger’s entire D2 column was recoded against the corrected rule before the freeze, while D1 and D3 verdicts were left untouched. This failure-and-fix is reported as a strength, not softened: it is the clearest evidence in the project that the reliability check was a real test, not a formality, catching a genuine rule gap rather than coder inattention, with the fix fully auditable in the diff between codebook versions.

Operating constraints. The instrument runs under seven numbered requirement rules (R-01 through R-07) that keep the coding auditable and replicable, covering source retrieval-dating, channel-taxonomy assignment, script-computed metrics, blinded reliability reporting, MRFM clause verifiability, no reproducible tamper or spoof technique, and dated-ADR logging of any post-freeze codebook change.

4 Results

The claim ledger (`ledger/claims.csv`) is frozen at v1.2, 125 rows, with all six degradation-condition columns (`d1_verdict`, `d2_verdict`, `d3_verdict` and their paraphrased notes) populated on every row. All figures in this section are read directly from `docs/headline-metrics.md`’s script-computed output (`python scripts/metrics.py`, run against the frozen 125-row ledger) or from a fresh, direct read of `ledger/claims.csv`’s `claim_type/d1_verdict` columns; per R-03, nothing here is computed or rounded by hand.

Headline metrics

- **M1 – CoT monopoly: 11.1%** ($n = 9$ T3/intent claims). About one in nine claims about *why* the agent did what it did rests solely on raw chain-of-thought (C1), with no corroborating channel at all.
- **M2 – Record survival under opacity: 90.4%** ($n = 125$ claims, overall). Split by claim type in Table 1 below.
- **M3 – Forgery exposure: 8.8%** ($n = 125$ claims). About one in eleven claims rests entirely on an agent-writable channel (C1/C2/C3) with no independent corroboration.

- **M4 – Assurance gap: 57.1%** ($n = 14$ T5/assurance claims). More than half of the record’s containment/completeness assurances are coded at verifiability tier V3 – not checkable by any third party via any stated procedure.
- **M5 – Broad channel fragility (supplementary): 78.4% / 47.2%** ($n = 125$ claims, both figures). 78.4% of all claims rest on one of the codebook’s four broadly fragile channels (C1/C2/C3/C6) as their primary channel; 47.2% – just under half of everything publicly known about this incident – rest on one of those channels with *no* corroboration of any kind.

M2 by claim type, and the figure

Table 1: M2 – record survival under D1 (opacity), by claim type. n and survival percentages are a direct read of `ledger/claims.csv` v1.2 (125 rows) and `match docs/headline-metrics.md`’s script output exactly.

Claim type		n	Survival under D1
T1	Event	58	100.0%
T2	Mechanism	42	90.5%
T3	Intent	9	22.2%
T4	Counterfactual	2	100.0%
T5	Assurance	14	92.9%

Figure 1 shows the same breakdown as claim counts rather than percentages, with each claim type’s bar segmented by D1 verdict (SURVIVES/DEGRADED/COLLAPSES). T1, T2, T4, and T5 are visibly SURVIVES-dominated – almost entirely green, with only a thin amber/red cap on T2 and T5. T3 is the one column that reads differently: green is a minority of the bar, and amber (DEGRADED) is the largest single segment. That visual asymmetry is M2’s finding in miniature – intent claims are the one claim type whose warrant is not, for the most part, independent of chain-of-thought being available.




Figure 1: Claim count by claim type (T1–T5), stacked by D1 (opacity) verdict. Green = SURVIVES, amber = DEGRADED, red = COLLAPSES. $n = 125$ claims total; per-type n as in Table 1. Generated by `scripts/make_results_figure.py` from `ledger/claims.csv`.

T3, two ways: strict survival versus any degradation

Table 1 reports T3’s strict D1-SURVIVES rate as 22.2% (2 of 9) – the same cut M2 uses for every other claim type. But collapsing DEGRADED and COLLAPSES into a single “T3 loses at least some precision without chain-of-thought” cut gives a different, arguably more informative number: **77.8% (7 of 9)** of T3 claims are DEGRADED or COLLAPSES under D1. Both numbers describe the same 9 rows and both are small- n – a

single row shifts either figure by roughly 11 percentage points, so neither should be read as precise. Read side by side, though, they tell a fuller story than either alone: the strict 22.2% survival figure by itself understates how CoT-dependent intent-attribution claims actually are, since it counts a claim as a full miss only if opacity destroys its warrant entirely; the 77.8% any-degradation figure captures every T3 claim that loses *some* precision the moment chain-of-thought is unavailable, which is the more complete picture of T3’s dependence on that one channel.

Reliability

Reliability was checked by a 25-row blind recode against Cohen’s κ per condition, against a 0.6 pass bar. D1 (opacity) came in at $\kappa = 0.7788$ and D3 (adversarial forgery) at $\kappa = 0.7794$, both comfortably above the bar and left untouched. D2 (unfaithfulness) came in at $\kappa = 0.5690$, below the bar, because the codebook’s shared corroboration logic did not account for D2’s failure mode being potentially systematic (an unfaithful self-report habit can recur consistently across a model’s transcripts, unlike a one-off forgery); the D2 rule was tightened and the full ledger’s D2 columns were recoded against it, logged as `adr/0001-d2-self-report-corroboration.md`. None of M1–M5 reads `d2_verdict` directly, so this fix changed the ledger’s reliability posture but not any number reported above.

The main threat to validity

Every SURVIVES verdict under D1 in this ledger is coded from what a source *says*, not from independently inspected evidence. A claim scores SURVIVES under opacity when its warrant would hold even without chain-of-thought – typically because a source asserts a second, corroborating channel (e.g., telemetry, an independent transcript, defender logs) – but this project has no way to inspect that asserted corroboration itself. If a cited corroborating channel is thinner, less independent, or less faithfully reported than the source claims, an unknown share of the 90.4% M2 headline figure is softer than it looks. This is the ceiling on M2 specifically, not a general disclaimer: M2 measures whether a claim’s stated warrant survives opacity, not whether that stated warrant was itself ever verified.

The 1w-01/1w-02 correction

1w-01 and 1w-02 are the record’s primary technical/mechanism sources, and were initially near-absent from the ledger after two retrieval attempts were blocked by a safety classifier. A narrower, targeted retry recovered 9 rows (C352–C360) at a category-only level of description, added to the ledger as a single disclosed exception to the P1.3 freeze (v1.1 → v1.2). All 9 rest solely on C5 (human testimony), a comparatively sturdy channel by the codebook’s own fragility ranking, with no corroboration. Adding them moved the fragility metrics in the *less* dramatic direction, not the more dramatic one: M3 moved from 9.5% (116 rows) to 8.8% (125 rows), and M5 moved from 84.5%/50.9% (116 rows) to 78.4%/47.2% (125 rows). What this means: the two sources whose absence was most plausibly suspected of understating this record’s fragility, once partially restored, pulled the fragility numbers *down* rather than up – evidence that the original 116-row figures were not inflated by omitting 1w-01/1w-02, and if anything were mildly conservative in the other direction. The correction is partial, not complete – the 9 recovered rows are coded at category-only description, coarser than the rest of the ledger – so this should be read as a directional finding, not a closed question; 1w-01/1w-02 content that a fuller retrieval might still surface could shift these numbers again in either direction. (A residual caveat about the coarseness of these 9 rows is restated as a standing limitation in Appendix A.)

5 MRFM v0.1

The Monitorability-Robust Forensic Minimum (MRFM) is this project’s proposed v0.1 minimum evidentiary standard for future agentic-AI incident reports: a set of clauses specifying what a lab, and where relevant its third-party compute providers, must retain, disclose, or make independently checkable so that a claim about an incident can be verified by an outside party without lab network access. Each of its 14 clauses is grounded directly in a specific, evidenced gap surfaced by this project’s 125-row claim ledger and its D1 (opacity) / D2 (unfaithfulness) / D3 (adversarial forgery) degradation analysis (Section 4), grouped below by the degradation each clause defends against, plus a fourth cross-cutting group addressing the verifiability of a lab’s own reporting process rather than of the agent’s record. Diffed against the two most directly relevant existing frameworks –

METR’s investigation-access list and GovAI’s incident-analysis retention list – **0 of the 14 clauses restate an existing requirement outright, 4 strengthen one by adding a verification mechanism it lacked (hash-chaining, signing, contemporaneous flagging, or per-claim procedure disclosure), and 10 are new relative to both sources.** This is best read as a finding about the state of existing guidance rather than a shortcoming of MRFM: METR’s and GovAI’s frameworks were written as general investigation-access and retention guidance, not against this specific incident, and did not anticipate several of the failure modes this incident’s ledger surfaced – most notably the entire D2 self-report-corroboration problem and three of the four cross-cutting process gaps, which sit outside what either framework was scoped to cover. Full six-field clause text (Requirement, Degradation defended, Evidence required, Verifier, Cost band, vs. prior art, plus each clause’s prior-art note) is in Appendix B; Table 2 below summarizes all 14 clauses.

Clause ID	Requirement (gist)	Degradation vs. prior art	
MRFM-D1-1	Retain targeted reasoning-trace excerpts for actions later classified as legitimacy/scope/objective judgments	D1	STRENGTHENS
MRFM-D1-2	Escrow a timestamped hash receipt for each retained excerpt with an independent party	D1	NEW
MRFM-D1-3	Pre-register retention-trigger criteria and a fixed disclosure-timing window before deployment	D1	NEW
MRFM-D2-1	Corroborate self-reported reasoning/intent claims with a non-self-report channel before treating them as fact	D2	NEW
MRFM-D2-2	Visibly label every reasoning/intent claim as self-reported or independently confirmed	D2	NEW
MRFM-D2-3	Disclose whether the implicated model has a measured CoT-faithfulness rate, or state that none exists	D2	NEW
MRFM-D3-1	Hash-chain agent-writable channels (C1–C3) at generation time in an append-only log	D3	STRENGTHENS
MRFM-D3-2	Cryptographically sign agent-writable entries, keyed to originating agent/session	D3	STRENGTHENS
MRFM-D3-3	Pre-commit hash-chain/signing verification material publicly before an evaluation begins	D3	NEW
MRFM-D3-4	Retain independently checkable network egress logs to cross-check agent-writable channels	D3	NEW
MRFM-CROSS-1	Require a stated, checkable verification procedure (or redaction summary) for every assurance claim	CROSS	STRENGTHENS
MRFM-CROSS-2	Maintain a dated, versioned public changelog for any report revision	CROSS	NEW
MRFM-CROSS-3	Publish quoted model output as machine-readable text alongside any figure/screenshot rendering it	CROSS	NEW
MRFM-CROSS-4	Obligate third-party compute providers to retain/disclose staging-environment logs	CROSS	NEW

Table 2: MRFM v0.1’s 14 clauses, condensed. Full requirement text, evidence required, verifier, cost band, and prior-art notes are in Appendix B.

6 Discussion

This sprint’s Track 2 judging criteria reward more than an accurate account of what happened: they also look for a causal account that generalizes past the single incident under study, one that says something testable about future cases rather than only this one. The finding this project is best positioned to offer against that bar is the T3 result. Intent-attribution claims are the one claim type this ledger’s own metrics single out as CoT-dependent – M1 shows 11.1% of T3 claims (1 of 9) rest solely on raw chain-of-thought with no corroborating channel at all, and M2’s per-claim-type breakdown shows T3 claims surviving the loss of chain-of-thought (D1, opacity) at only 22.2%, against 89–100% survival for every other claim type (Section 4). That gap is not an idiosyncrasy of this incident, this lab, or this model. It follows from a structural asymmetry in what different claim types assert. An event claim (T1) or a mechanism claim (T2) asserts something about the world external to the acting system – a timestamp, a file touched, an API called – and such assertions can, in principle, leave independent traces in channels the agent itself did not write: host-side logs, cloud-provider audit trails, a second party’s own record. An intent claim (T3) asserts something about the acting system’s internal state – what it judged, and why it chose one action over another – and absent a disclosed, faithful chain-of-thought, internal state leaves no such external trace to find. It can only be inferred behaviorally, from what the agent did, and behavioral inference is a

strictly lower-fidelity signal than a contemporaneous, verbalized account of the reasoning itself would be. This is a prediction, not just a description: any future agentic incident, regardless of which lab or which model produced it, should show the same shape – event and mechanism claims surviving the loss of chain-of-thought at rates near this ledger’s T1/T2 figures, intent claims collapsing toward its T3 figure – because the asymmetry is a fact about what an intent claim structurally is, not a fact about this incident’s particular evidentiary luck. That is the sense in which this project’s sharpest finding is not only an audit of one incident’s public record, but a testable claim about the next one.

This structural asymmetry has a direct corollary once chain-of-thought itself stops being treated as a fixed feature of future models. The CoT-monitorability-fragility literature already surveyed in Related Work (Section 2) argues that today’s externalized, human-readable reasoning trace is an emergent property of current training methods, not a guarantee: Korbak et al.’s cross-lab position paper names a shift toward less monitorable, less verbalized reasoning – optimizing directly against a CoT monitor, or moving toward non-linguistic latent reasoning – as a live risk, and Baker et al. show a version of this already happening under direct optimization pressure, where a model taught to evade a CoT monitor keeps misbehaving while its trace simply stops describing why. Chen et al.’s hint-based faithfulness tests already put verbalization rates well below 50% in some settings even without adversarial pressure on the monitor. If the trend these sources describe continues, the T3/M1 dependency finding above does not shrink – it grows. The one channel intent-attribution claims currently lean on hardest, imperfectly, is also the channel this trend degrades fastest, while the alternative – behavioral inference from actions alone – gains no corresponding improvement as a window into internal state. A world with less monitorable, less verbalized model reasoning is a world where this ledger’s 22.2% D1-survival rate for T3 claims looks less like a floor for future incidents’ intent claims and more like a ceiling.

The eight RFI questions in `docs/rfi-questions.md` exist so that this same channel-fragility finding is not only a description of a problem but something a regulator – or a Track 3 team building governance, verification, or enforcement tooling on top of this sprint’s output – can act on directly. Each question is anchored to the specific `claim_ids` where the ledger itself, via a `verifiability=V3` coding or a `COLLAPSES` verdict under D1, D2, or D3, already flags the public record as weakest, rather than to a fresh judgment call made for the RFI document itself; and each carries a stated “what answer would settle this” and “what answer would NOT settle this” pair, so a recipient’s response can be checked against a fixed bar instead of read for tone or reassurance value. That is exactly the shape of question a regulator needs in hand to actually send, rather than a general request for more transparency, and exactly the kind of ready-made, evidence-anchored artifact a Track 3 team could pick up and operationalize without re-deriving it from the raw 125-row ledger – a concrete, checkable question with a stated pass/fail bar, which is the other half of what this track’s judging criteria reward alongside the causal claim above.

A month of follow-up would strengthen the method itself, not just this one application of it. Three additions matter most. First, multi-rater coding throughout the ledger, rather than the single blind-recode subsample this project used to compute Cohen’s κ per degradation condition (Section 4; `docs/reliability-report.md`) – a full second coder on every row would tighten every headline number’s confidence interval, not just confirm that the codebook survived one spot-check. Second, running the same instrument – codebook, degradation conditions, and headline metrics unchanged – against a second, independent agentic incident, to test whether the channel-fragility pattern found here, and the T3-specific prediction above in particular, replicates outside this one case, or whether this incident’s specific mix of disclosed and withheld evidence was unusual. Third, piloting the RFI questions in `docs/rfi-questions.md` against an actual regulatory review process, or at minimum against feedback from someone who has drafted or received a real RFI, to check whether the questions as phrased function as intended once someone outside this project tries to use them. (MRFM v0.1’s own limitations and its own month-of-follow-up list, in `mrfrf/mrfrf-v0.1.md`, are scoped to the fourteen-clause table specifically and are not repeated here.)

Returning to this project’s thesis: the best-powered evidence for it is not the T3 finding this section leads with structurally, but M5. Across all 125 claims in the ledger, 78.4% rest on a primary channel the codebook itself ranks as fragile (C1, C2, C3, or C6), and 47.2% – just under half of everything publicly known about this incident – rest on such a channel with no corroborating channel at all. M1’s 11.1% ($n=9$) is the sharpest, most viscerally specific illustration of the same underlying pattern, and the causal argument above is where that pattern’s mechanism is easiest to state precisely, but M5 is the number that should carry this paper’s statistical weight when the claim is about the incident’s public record as a whole: not a fragile statistic about nine claims, but a well-powered one about all of them.

In sum: this project set out to measure, rather than assume, how much of the public understanding of one agentic-AI incident rests on evidentiary channels likely to be absent, unfaithful, or forged next time. The answer is a large fraction under any of the cuts this ledger supports; the mechanism behind the sharpest slice of that

fraction, intent attribution, is structural rather than incidental to this incident; and the direction of travel in how future models reason makes that mechanism more consequential, not less. MRFM v0.1 (Section 5) and the RFI questions above are this project’s answer to what breaks next and what to do before it does; whether they hold up is exactly the kind of check the follow-up work above is designed to run.

References

1. Cohen, J. (1960). “A coefficient of agreement for nominal scales.” *Educational and Psychological Measurement*, 20(1), 37–46.
2. Korbak, T., Balesni, M., Barnes, E., et al. (2025). “Chain of Thought Monitorability: A New and Fragile Opportunity for AI Safety.” Cross-lab position paper. <https://arxiv.org/abs/2507.11473>
3. Baker, B., Huizinga, J., Gao, L., Dou, Z., Guan, M. Y., Madry, A., Zaremba, W., Pachocki, J., & Farhi, D. (2025). “Monitoring Reasoning Models for Misbehavior and the Risks of Promoting Obfuscation.” OpenAI. <https://arxiv.org/abs/2503.11926>
4. Chen, Y., Benton, J., Radhakrishnan, A., Uesato, J., Denison, C., Schulman, J., Somani, A., et al. (2025). “Reasoning Models Don’t Always Say What They Think.” Anthropic. <https://arxiv.org/abs/2505.05410>
5. Lanham, T., Chen, A., Radhakrishnan, A., Steiner, B., Denison, C., Hernandez, D., Li, D., Durmus, E., Hubinger, E., Kernion, J., et al. (2023). “Measuring Faithfulness in Chain-of-Thought Reasoning.” Anthropic. <https://arxiv.org/abs/2307.13702>
6. Turpin, M., Michael, J., Perez, E., & Bowman, S. R. (2023). “Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting.” *NeurIPS 2023*. <https://arxiv.org/abs/2305.04388>
7. Arcuschin, I., Janiak, J., Krzyzanowski, R., Rajamanoharan, S., Nanda, N., & Conmy, A. (2025). “Chain-of-Thought Reasoning In The Wild Is Not Always Faithful.” <https://arxiv.org/abs/2503.08679>
8. Coalition for Secure AI (CoSAI), OASIS Open, Workstream 2 (2025). “AI Incident Response Framework V1.0.” Released November 18, 2025. <https://github.com/cosai-oasis/ws2-defenders/blob/main/incident-response/AI-Incident-Response.md>
9. Nelson, A., Rekhi, S., Souppaya, M. (NIST), & Scarfone, K. (Scarfone Cybersecurity) (2025). “NIST SP 800-61r3 – Incident Response Recommendations and Considerations for Cybersecurity Risk Management: A CSF 2.0 Community Profile.” Final release April 2025. NIST. <https://csrc.nist.gov/pubs/sp/800/61/r3/final> (DOI: 10.6028/NIST.SP.800-61r3)
10. Cloud Security Alliance (CSA) (2026). “AI Controls Matrix v1.1 – AIS-13 (AI Sandboxing).” July 2026. <https://cloudsecurityalliance.org/artifacts/ai-controls-matrix-v1-1>
11. Ezell, C., Roberts-Gaal, X., & Chan, A. (Centre for the Governance of AI) (2025). “Incident Analysis for AI Agents.” Proceedings of AIES 2025; arXiv posted August 19, 2025. <https://arxiv.org/abs/2508.14231>
12. METR (2026). “How independent researchers could investigate AI propensities after misalignment incidents.” July 28, 2026. <https://metr.org/blog/2026-07-28-investigating-ai-propensities-after-incidents/>
13. Centre for Long-Term Resilience (CLTR) (2026). “The Loss of Control Observatory: A Prototype to Detect Real-World AI Control Incidents.” February 2, 2026. <https://www.longtermresilience.org/reports/the-loss-of-control-observatory-a-prototype-to-detect-real-world-ai-control-incidents/>
14. Frontier Model Forum (2026). “Information Sharing, Incident Reporting, and Incident Response for Frontier AI Risks.” Issue brief, May 12, 2026. <https://www.frontiermodelforum.org/issue-briefs/information-sharing-incident-reporting-and-incident-response-for-frontier-ai-risks/>
15. Dixon, R. B. L., & Frase, H. (CSET, Georgetown University) (2025). “AI Incidents: Key Components for a Mandatory Reporting Regime.” January 2025. <https://cset.georgetown.edu/publication/ai-incidents-key-components-for-a-mandatory-reporting-regime/>

16. Wei, K., & Heim, L. (2026). “Designing Incident Reporting Systems for Harms from General-Purpose AI.” *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 40, AAAI-26; arXiv v1 November 8, 2025, revised April 2026. <https://arxiv.org/abs/2511.05914>

A Limitations and Dual-Use Considerations

Limitations

This audit’s conclusions are bounded by six limitations, each of which would change how the headline metrics should be read if ignored.

Single-coder ledger, blinded re-code only. The 125-row claim ledger was built and coded by a single researcher, not by independent multi-rater coding across the whole ledger. The reliability check that exists is a 25-row blind re-code (Rater A’s original P1.1 verdicts held out and compared against a cold, independently re-coded pass over the same 25 rows), scored with Cohen’s κ per degradation condition. That check found $\kappa = 0.7788$ (D1, opacity) and $\kappa = 0.7794$ (D3, adversarial forgery), both above the 0.6 bar, but $\kappa = 0.5690$ on D2 (unfaithfulness) – below it. The D2 codebook rule was subsequently fixed and the whole ledger’s D2 column was recoded against the corrected rule before freeze. This resolves the specific rule ambiguity the recode surfaced; it does not substitute for full independent multi-rater coding across all 125 rows, and the remaining 100 rows outside the 25-row subsample were never independently re-coded by a second rater under either the original or the fixed rule.

Corroboration taken at face value. Claims are coded as *reported*. A claim can be scored SURVIVES under D1 (opacity) because a source *asserts* the existence of corroborating telemetry, logs, or a second channel that this project has no independent means to inspect. This project cannot verify that an asserted corroborating channel actually exists, was actually consulted, or would actually hold up if produced. This is the ceiling on M2 (record survival under opacity, 90.4% overall) in particular: a claim scored as surviving opacity is a claim whose *stated* evidentiary chain does not require chain-of-thought, not a claim this project has confirmed would survive an actual forensic examination. We restate this here as a named limitation on the audit’s method, not only as a caveat embedded in the Results narrative.

Channel-biased claim population. The claim population is drawn from what the public record chose to publish – incident reports, blog posts, and third-party investigations that already made editorial choices about which evidence to disclose, quote, or summarize. Channels that were never written up publicly at all cannot appear in this ledger, and channels that labs chose to describe only at a category level (rather than quoting directly) are underrepresented in the ledger’s degree of technical detail regardless of how load-bearing they actually were. The ledger therefore measures channel dependence *within the published record*, not within the full evidentiary record the investigating institutions actually held.

Degradation verdicts are analytic, not empirical. The D1/D2/D3 verdicts (SURVIVES / DEGRADED / COLLAPSES) are judgments about whether a claim’s stated evidentiary warrant would hold up under a stress condition – opacity, unfaithfulness, or adversarial forgery – reasoned from the claim and its cited channel(s) as reported. They are not empirical measurements of what an investigator with real access to raw logs, model weights, or infrastructure could actually recover, reconstruct, or independently verify under any of these conditions. No degradation condition in this project was tested against real systems; all three are applied as analytic overlays on the claim ledger.

Single-incident calibration. The codebook, the three degradation conditions, and the five headline metrics were all built against one incident. Nothing in this project establishes that the same channel codes, verifiability tiers, or degradation rules would classify claims from a different incident the same way, or that the specific fragile-channel finding (M5: 78.4% of all 125 claims rest on a codebook-fragile primary channel) generalizes beyond this one case. The instrument’s transfer to other incidents is untested.

Small-N metrics. Two of the headline results rest on thin denominators and should be read as directional rather than precise. M1 (CoT monopoly, 11.1%) is computed over only 9 T3 (intent) claims in the entire ledger; a single row moving to or from the sole-C1 condition would shift the reported figure by roughly 11 percentage points. M2’s per-claim-type breakdown for T4 (counterfactual claims, 100% survival under D1) rests on only 2 claims total – two data points cannot confirm or refute the sprint plan’s own prediction that T3 and T4 claims would both fail to survive opacity; the honest reading is that T4 did not collapse in this ledger’s two counterfactual claims, not that the plan’s prediction about T4 was confirmed. Both caveats are documented in `docs/headline-metrics.md` and are restated here as standing limitations on how much confidence M1 and the T4 slice of M2 can bear.

Residual 1w-01/1w-02 category-level gap. The ledger was frozen once at v1.1 (116 rows) and reopened a single, disclosed time to add 9 rows (C352–C360) recovered from sources 1w-01 and 1w-02, bringing it to v1.2 (125 rows). That recovery pulled M3 and M5 toward less dramatic values than the pre-recovery figures (M3: 9.5% → 8.8%; M5: 84.5%/50.9% → 78.4%/47.2%), which is itself evidence the original numbers were not inflated by the omission in the other direction. This is a separate and narrower point from that directional finding, however: the 9 recovered rows were extracted and coded at a category-only level of description – section headings and topic identification only, never the underlying technical content, commands, or payloads themselves. The fuller technical detail present in 1w-01/1w-02’s own source material is still not reflected anywhere in the ledger. What this project reports about those two sources’ contribution to the record is therefore a coarser summary than what the sources themselves actually contain, and that gap should not be read as closed by the 9-row addition – only partially narrowed.

Dual-Use Considerations

This audit names which evidentiary channels are load-bearing for the public’s understanding of this incident. That is also, unavoidably, a map of which channels an adversary – or an institution with an incentive to control the narrative of a future incident – would most want to suppress, withhold, or forge next time. We take this risk seriously and adopt two mitigations, both already scoped in the sprint plan before this project’s data collection began.

First, the analysis is kept at the level of *channel classes* (raw chain-of-thought, inter-agent message content, tool-call/action traces, lab self-assertion, and so on) rather than at the level of any specific, reproducible tamper or spoofing technique. Per rule R-06, no note field in the ledger, no discussion in this report, and no supporting document in this repository reproduces a forgery or tampering mechanism at a level of detail that would let a reader carry it out. Where the public source material documents a successful tool-call-spoofing technique used during the incident, this project logs its existence and category only – consistent with how every D3 (adversarial forgery) note in the ledger is written, and with how the excerpt inventories treat comparable technique-level content throughout. We state explicitly, per the sprint plan’s own instruction on this point: this project declines to reproduce the documented spoofing technique, even though it is described in a public source available to us. Publishing a channel-class ranking is a different act from publishing a recipe, and this project confines itself to the former.

Second, the constructive half of this project – the v0.1 minimum forensic standard (MRFM) – is written so that its defensive value exceeds its offensive value. Its clauses tell a defending institution which channels need redundancy, retention, and attestation; they do not tell an attacker anything about specific channels’ weaknesses that could not already be inferred from the same published incident reports this audit draws on. A channel ranking derived entirely from public retrospectives is not new information handed to a capable adversary – the labs’ own reports already disclose which channels existed and how they behaved under stress. What this project adds is a systematic account of which channels the *public record’s own claims* depend on, framed toward the question of what an institution should preserve, corroborate, and attest to before the next incident, not toward what a future intruder should target.

B Full MRFM v0.1 Clause Table

This appendix reproduces, in full, all 14 clauses of MRFM v0.1 – the Monitorability-Robust Forensic Minimum introduced in Section 5 – grouped by the degradation condition each clause defends against (D1 opacity, D2 unfaithfulness, D3 adversarial forgery) plus a fourth group of cross-cutting clauses about the verifiability of a

lab’s own reporting process rather than of the agent’s record. Each clause is given in its full six-field form (Requirement, Degradation defended, Evidence required, Verifier, Cost band, vs. prior art) together with its prior-art note, followed by the diff summary against METR’s and GovAI’s frameworks and the standard’s stated limitations. This is the frozen **v0.1** — **frozen 2026-09-12** version of `mrfm/mrfm-v0.1.md`; nothing below has been added to, trimmed from, or reworded relative to that source beyond Markdown-to-L^AT_EX formatting.

Group D1: Opacity

Chain-of-thought or reasoning traces are never logged, or are logged but never disclosed to investigators (`docs/codebook.md` §2.1, channel C1). `docs/headline-metrics.md` M2 shows this is not a uniform problem: only T3 (intent) claims collapse under D1 (22.2% survival); the other four claim types survive at 89%+. The three clauses below are scoped to that specific, evidenced gap rather than proposing blanket CoT retention.

MRFM-D1-1: Targeted retention of legitimacy/scope/objective reasoning traces

Requirement

For any agent action or decision that is later classified — under the ex-ante criteria required by MRFM-D1-3 — as involving a judgment about the legitimacy, scope, or objective of a task (the category of content that generates T3-type claims per `docs/codebook.md` §2.2/§3.4), the lab must retain the corresponding reasoning-trace segment and make it available to a credentialed incident investigator on request. This is a targeted duty, not a blanket mandate to retain or disclose all chain-of-thought: `docs/headline-metrics.md` M2 shows T1/T2/T4/T5 claims already survive D1 at 89%+ from other channels, so a blanket retention rule would impose cost (`docs/prior-art.md`’s CoSAI entry notes finite storage/compute means retention objectives should be set deliberately, not by assuming unlimited logging) without addressing the one claim type that actually depends on it.

Degradation defended

D1 (opacity)

Evidence required

A retained, timestamped reasoning-trace excerpt keyed to the specific flagged action/decision, stored separably from bulk CoT logs (so it can be produced on its own rather than requiring full-transcript disclosure), with a record showing it was flagged and retained contemporaneously with the action rather than reconstructed after an incident became public.

Verifier (no lab network access needed)

A third-party investigator (e.g. a body in the METR/AISI mold, per `docs/prior-art.md`’s METR entry) checks the excerpt’s timestamp and content against the escrowed hash receipt required by MRFM-D1-2, and against other independently-available contemporaneous records (e.g. a C3 tool-call timeline), once the excerpt is produced — the check itself needs only the produced artifacts, not the lab’s live systems.

Cost band Medium — most labs already generate CoT internally; the added cost is a tagging/retrieval pipeline that flags and separably stores excerpts matching the MRFM-D1-3 criteria and a disclosure process for producing them, roughly 1-4 engineer-weeks rather than new capture infrastructure.

vs. prior art

STRENGTHENS

Prior-art note

GovAI’s “adjacent material” item already asks developers to retain “records of the agent’s reasoning/decision process,” and METR’s “Full transcripts / reproducible environments” item asks for complete transcripts sufficient to reconstruct an incident; MRFM-D1-1 narrows this to a criteria-gated subset and adds contemporaneous flagging plus separable storage that neither source specifies.

MRFM-D1-2: Third-party escrow receipt for retained reasoning traces

Requirement

At the time a reasoning-trace excerpt is flagged and retained under MRFM-D1-1, the lab must compute a cryptographic hash (or equivalent commitment) of that excerpt and deposit the hash plus a timestamp with an independent party not under the lab’s operational control — e.g. a third-party evaluator such as METR or a national AI safety institute, or a dedicated escrow service ([docs/prior-art.md](#)’s METR entry already documents this kind of body reviewing an incident involving this same lab pairing, so the relationship type is not novel). This exists because a retention duty alone is only checkable by asking the lab to attest that it complied — a receipt held by an independent party lets an outsider confirm retention occurred without relying on the lab’s word.

Degradation defended

D1 (opacity)

Evidence required

A timestamped hash receipt held by the escrow party, generated at or near the time of the underlying action rather than backfilled once an incident is already public, covering the specific excerpt(s) retained under MRFM-D1-1.

Verifier (no lab network access needed)

Any investigator requests the receipt directly from the escrow party — not the lab — and, once the lab later discloses the underlying trace content (via MRFM-D1-1 or a subsequent incident report), independently recomputes the hash from the disclosed content and checks it against the escrowed receipt. Both steps use only the escrow party’s records and the disclosed content; no access to the lab’s internal network is needed.

Cost band High — requires a standing relationship with an independent party and ongoing operational integration into the evaluation/deployment pipeline, a genuinely recurring cost rather than a one-time build.

vs. prior art

NEW

Prior-art note

Neither METR’s investigation-access list nor GovAI’s retention list mentions cryptographic hashing or third-party escrow of retained evidence — both stop at “retain it and grant investigator access,” with no mechanism for proving retained content wasn’t altered or backfilled after the fact.

MRFM-D1-3: Pre-registered retention-trigger criteria and disclosure window

Requirement

Before deploying or evaluating an agent capable of extended autonomous operation, a lab must publish the concrete criteria it will use to classify an action or decision as “involving a judgment about legitimacy, scope, or objective” (the MRFM-D1-1 trigger) and commit to a fixed disclosure window (e.g., a stated number of days) measured from the point an event is formally classified as an incident. Both the criteria and the timing commitment must be fixed in advance, not decided once an incident is already underway — otherwise MRFM-D1-1’s duty is gameable by narrowing what counts as qualifying content, or by indefinitely deferring disclosure, only after a lab’s incentives have changed. [docs/prior-art.md](#)’s CoSAI entry and NIST SP 800-61r3 entry both frame retention and response commitments as things that should be set at design time, before an incident occurs, which is the same principle applied here to the classification trigger itself rather than just storage volume.

Degradation defended

D1 (opacity)

Evidence required

A dated policy document, published or otherwise made available to credentialed investigators prior

to the deployment/evaluation run in question, stating the classification criteria and the disclosure-timing commitment; and, for any actual incident, the real disclosure date to compare against that commitment.

Verifier (no lab network access needed)

Any third party (investigator, regulator, journalist) compares the policy document’s publication date against the deployment’s start date, and separately compares the actual disclosure date for a given incident against the policy’s stated window. Both checks require only two dated artifacts already outside the lab’s systems — the published policy and the disclosed incident material — not lab network access.

Cost band Low — a policy/documentation exercise defining criteria and a timing commitment in advance; it adds no new capture infrastructure beyond what MRFM-D1-1 already requires, so it is achievable as a policy change in under an engineer-week.

vs. prior art

NEW

Prior-art note

Neither METR nor GovAI address publishing classification criteria or a disclosure-timing commitment in advance of deployment; both lists specify what to retain, not when the rules governing retention/disclosure must be fixed relative to an incident (this clause’s design-time framing instead draws on the CoSAI/NIST entries already cited in its own text, which sit outside the METR/GovAI scope this diff targets).

Group D2: Unfaithfulness

D2 (`docs/codebook.md` §2.3) is about whether a *disclosed* reasoning trace faithfully reports the computation that produced the action — distinct from D1, which is about whether the trace exists or is disclosed at all. A lab can satisfy every D1-facing clause above and still produce an incident report that over-trusts an unfaithful self-report, which is the gap these three clauses target, drawn directly from a reliability failure in this project’s own coding process (`adr/0001-d2-self-report-corroboration.md`).

MRFM-D2-1: Independent-corroboration requirement for self-reported reasoning

Requirement

Before an incident report presents a claim about the acting agent’s stated reasoning, belief, or intent as established fact, the report must corroborate that claim with at least one channel that (a) is not itself another instance of the same kind of agent self-report — not a second transcript, a second agent instance, or an aggregate/sampled population restating the same kind of self-narrated content — and (b) independently establishes the *same specific scope* the claim asserts, not merely an adjacent broader or narrower fact. Where a claim bundles a self-report sub-assertion (a stated reason, belief, or intent) together with a technical/behavioral sub-assertion that genuinely is independently corroborated, the report must state its conclusion against the claim’s most specific stated content, not a softened composite that borrows credibility from the corroborated half to cover the uncorroborated half.

Degradation defended

D2 (unfaithfulness) — directly generalizes the finding in `adr/0001-d2-self-report-corroboration.md` and `docs/codebook.md` §3.7: an unfaithful self-report can be a systematic property of how a model narrates itself, so a second self-report-type channel does not test for faithfulness the way it would test for mere availability (D1) or one-off forgery (D3).

Evidence required

For each reasoning/intent claim the report makes, a stated mapping to the specific non-self-report channel that corroborates it (e.g., defender-side telemetry, a raw tool-call/action-log entry distinct from the agent’s own narrated reason for that action, human testimony, or an investigator-verified record) and confirmation that the channel’s content matches the claim’s specific asserted scope.

Where no such channel exists, the report must say so explicitly rather than let the self-report stand uncontested as fact.

Verifier (no lab network access needed)

Any external reader of the published report — a journalist, academic auditor, or affected third party — by checking the report’s own text: does each reasoning/intent claim cite a channel other than a second self-report, and does the cited channel’s described content actually cover the claim’s specific scope rather than an adjacent one? This is a document-literacy check against the report as published; it requires no transcript access, no technical infrastructure, and no lab cooperation.

Cost band Low — under an engineer-week. The corroborating evidence, where it exists, is typically already collected during the lab’s own investigation; the cost is the report-writing discipline of mapping each reasoning/intent claim to its specific corroborating channel (or explicitly flagging its absence) rather than generating new evidence.

vs. prior art

NEW

Prior-art note

METR’s list asks that investigators be given transcripts, interviews, and reproduction access so they can independently check things themselves; GovAI’s retention list is silent on report-writing epistemics entirely. Neither requires the report’s own authors to corroborate a self-reported reasoning claim with a non-self-report channel before presenting it as established fact.

MRFM-D2-2: Self-report vs. independently-confirmed labeling requirement

Requirement

Every sentence or claim in an incident report that asserts a reasoning-, intent-, or belief-type fact about the agent must be visibly labeled, in the report’s own text or an accompanying claim table, as either “self-reported” (its sole basis is the agent’s own narrated output, however many times repeated) or “independently confirmed” (it satisfies MRFM-D2-1’s corroboration test). A claim that bundles a self-reported sub-assertion with an independently-confirmed sub-assertion must be split into its components for labeling purposes rather than given one blended label that implies more confirmation than the self-reported portion actually has.

Degradation defended

D2 (unfaithfulness) — makes the self-report/corroboration distinction that MRFM-D2-1 requires labs to *apply* internally also visible externally, so a reader is not left to infer from prose tone alone whether “the agent explained that it acted because...” is a confirmed fact or an unverified self-narration.

Evidence required

The published incident report itself, containing per-claim or per-paragraph provenance labels, or a claim-provenance table attached to or embedded in the report, consistent with the corroboration mapping required by MRFM-D2-1.

Verifier (no lab network access needed)

Any member of the public reading the released report. The check is purely textual: are the labels present, and are they internally consistent with the evidence the report itself cites for each claim? No special technical background, transcript access, or lab cooperation is required — this is the same zero-network-access bar as MRFM-D2-1’s verification.

Cost band Low — under an engineer-week. This is a drafting/formatting convention applied at report-writing time, comparable in effort to standard citation formatting, not a new technical capability.

vs. prior art

NEW

Prior-art note

No labeling/provenance convention for individual claims appears in either GovAI’s retention list or METR’s investigation-access list; both are silent on how a published report should visibly distinguish self-reported from independently-confirmed content.

MRFM-D2-3: Faithfulness-measurement disclosure requirement

Requirement

If an incident report’s narrative relies on any disclosed reasoning trace (channel C1) from a given model or model version, the report must state whether that model has been subject to an established faithfulness-measurement procedure of the kind documented in the public CoT-faithfulness literature (e.g., perturbation- or hint-based tests that check whether a model’s stated reasoning changes consistently with its behavior) and, if so, disclose the measured faithfulness/verbalization rate and the conditions under which it was measured. If no such measurement exists for the model in question, the report must say so explicitly rather than let the narrative implicitly assume the disclosed CoT is a reliable report of the computation that produced the action.

Degradation defended

D2 (unfaithfulness) — generalizes the same ADR 0001 insight one level up from individual claims: because an unfaithful narration habit can be a systematic property of a model (not a one-off slip a single corroborating channel would catch), a reader needs to know the model’s *baseline* measured faithfulness rate, not just whether any one claim happens to have a corroborating channel. This is grounded in the published finding that verbalization/fidelity rates in reasoning models have been measured well below 50% in some evaluated settings, with a documented trend toward less faithful reasoning as models scale on some tasks (see `docs/prior-art.md` §1, citing the Anthropic and Turpin et al. faithfulness literature) — monitorability of this kind is an emergent property that a report should not silently assume persists, per the position paper on CoT monitorability’s fragility also cited there.

Evidence required

A cited faithfulness-measurement result for the implicated model — methodology used and date measured — or, absent one, an explicit disclosed statement that no such measurement has been performed for that model.

Verifier (no lab network access needed)

Any external reviewer with public familiarity with the CoT-faithfulness literature (an academic reviewer, an AI-safety-focused journalist, or a third-party auditor) can check whether the report cites a specific, dated faithfulness measurement, and — where the model family has been studied in published faithfulness research — can check the cited figure for plausibility against that public literature. No lab-internal access is needed to perform this check; it only requires the published report and publicly available faithfulness studies.

Cost band Medium — 1-4 engineer-weeks for a lab with no existing faithfulness-testing pipeline to run and report a baseline measurement using established methodology adapted to the implicated model. Labs with an existing internal faithfulness-testing practice incur only the low cost of disclosure; the band is set at Medium because the requirement must still be satisfiable by a lab starting from zero.

vs. prior art

NEW

Prior-art note

Neither METR nor GovAI mention faithfulness-measurement disclosure at all; this clause is grounded instead in the CoT-faithfulness position-paper literature (`docs/prior-art.md` §1) — a different sweep than the METR/GovAI investigation-and-retention sweep this diff targets.

Group D3: Adversarial forgery

Agent-writable channels (C1 raw chain-of-thought, C2 inter-agent message content, C3 tool-call/action traces) may be forged, and the incident record documents this as attempted and, in one technique, successful (`docs/codebook.md`

§2.3). `docs/headline-metrics.md` reports M3 (forgery exposure) at 8.8% of all 125 ledger claims resting solely on an agent-writable channel with no corroboration, and the supplementary M5 shows just under half of everything publicly known about the incident rests on a channel this project’s codebook already flags as forgeable, with nothing else backing it up.

MRFM-D3-1: Hash-chained, append-only logging of agent-writable channels

Requirement

Every channel the codebook classifies as agent-writable (C1 raw CoT, C2 inter-agent message content, C3 tool-call/action traces) must be recorded to an append-only log at generation time, where each entry’s stored hash is computed over that entry’s content *and* the immediately preceding entry’s hash. No mechanism may exist (or, if one exists for operational reasons, its use must itself be logged in an equally tamper-evident way) for altering or removing an already-written entry without breaking the hash chain from that point forward.

Degradation defended

D3 (adversarial forgery).

Evidence required

The complete hash-chained log for the incident window, plus the generation-time hash of the final entry recorded at (or as close as practicable to) the time it was written — e.g., via a timestamped export, a periodic checkpoint published outside the lab’s own control, or a commitment recorded in a public transparency log per MRFM-D3-3. A lab claiming this control is in place must be able to produce the full chain on request, not just assert its existence (avoiding the V3 “assurance with no stated procedure” failure pattern `docs/headline-metrics.md` documents for M4).

Verifier (no lab network access needed)

Any third party holding a full or partial export of the log can independently recompute the hash chain and confirm it is unbroken, using only the disclosed data and the hashing algorithm’s public specification — no access to the lab’s internal systems is required. A broken or non-recomputable chain is itself the finding: it proves *that* something was altered after the fact, without the verifier needing to know how any specific alteration technique works.

Cost band Low — hash-chaining (each entry storing a hash of itself plus the prior entry’s hash) is a well-understood, computationally cheap logging pattern with no dependency on key management or external infrastructure; the main cost is disciplined integration into existing logging pipelines for C1/C2/C3, not new cryptographic machinery.

vs. prior art

STRENGTHENS

Prior-art note

GovAI’s “Activity logs” item already asks for “a complete, timestamped record of the agent’s actions and outputs/decisions”; MRFM-D3-1 adds the hash-chaining mechanism (each entry’s hash covers its own content plus the prior entry’s hash) that makes tampering with an already-written entry detectable — a verification property GovAI’s plain “complete record” ask does not specify.

MRFM-D3-2: Cryptographic signing of agent-writable channel entries, keyed to originating agent/session

Requirement

Every entry written to an agent-writable channel (C1, C2, C3) must be signed with a key cryptographically tied to the specific agent instance and session that produced it, such that authorship of any given entry — which agent, in which run, wrote this — is independently verifiable from the signature rather than resting on an unauthenticated channel or on the channel’s own self-reported attribution.

Degradation defended

D3 (adversarial forgery).

Evidence required

The signed entries themselves, the key-issuance record mapping each agent/session identifier to its key (or key fingerprint), and — for the specific gap this clause targets — a record of which handles claimed which fingerprints, in the shape METR’s own PDF release already demonstrates is producible (a structured signed-message identity table, not a new invention). A lab must disclose the signing scheme’s verification method (e.g., which public-key algorithm, and how fingerprints map to sessions) at or before evaluation start, not after an incident occurs.

Verifier (no lab network access needed)

Any third party holding the disclosed public keys (or fingerprints) and the signed entries can independently verify each entry’s signature against the claimed originating agent/session, without needing lab cooperation at verification time. An entry whose signature does not verify against its claimed origin is evidence of forged or misattributed authorship, without the verifier needing to know how forgery of the *unsigned* channel would have worked.

Cost band Medium — requires key-generation and key-issuance infrastructure scoped to every agent instance/session (higher setup cost than MRFM-D3-1’s hashing), plus a disclosure process for the verification material, but reuses standard public-key signing rather than novel cryptography; the incident record itself shows agents improvising a version of this scheme under far worse conditions (no prior key infrastructure, no lab support), so a lab-provisioned version should not exceed that in difficulty.

vs. prior art

STRENGTHENS

Prior-art note

Same GovAI “Activity logs” baseline as MRFM-D3-1; this clause adds cryptographic signing keyed to originating agent/session so authorship is independently verifiable, rather than resting on the channel’s own unauthenticated, self-reported attribution — a mechanism GovAI’s record-keeping ask does not specify.

MRFM-D3-3: Pre-committed, publicly checkable verification material**Requirement**

At or before the start of an evaluation, a lab must publish (or commit to, via a public transparency log or equivalent third-party-witnessed mechanism) the verification material needed to check MRFM-D3-1 and MRFM-D3-2’s integrity guarantees — the hash-chain’s starting/checkpoint state and the signing scheme’s public keys or fingerprints — such that this material is fixed and independently datable *before* any incident this evaluation might produce, and cannot be substituted or backdated after the fact.

Degradation defended

D3 (adversarial forgery) — specifically, the case where a lab could otherwise generate a hash-chain or signing keys *after* an incident and claim they were always in place.

Evidence required

A dated, independently-witnessed publication record (e.g., a transparency-log inclusion proof, a timestamped public post, or an equivalent third-party-verifiable commitment) of the verification material, made before the evaluation window begins.

Verifier (no lab network access needed)

Any third party can check the publication timestamp and inclusion proof against the transparency log or publication venue directly — this is the mechanism that makes MRFM-D3-1/D3-2 checkable using only the published key/log and the disclosed data, with no lab network access or cooperation needed at verification time.

Cost band Low — publishing a commitment (a hash, a key, or a fingerprint) to an existing public transparency-log service or a simple timestamped public post is low-effort and does not require building new infrastructure beyond what MRFM-D3-1/D3-2 already require the lab to generate.

vs. prior art
NEW

Prior-art note

Neither source addresses pre-committing or publicly timestamping verification material (hash-chain checkpoints, signing keys) before an evaluation begins; this closes a backdating loophole that GovAI's and METR's lists don't anticipate at all.

MRFM-D3-4: Retained, independently checkable egress logs

Requirement

Network egress logs covering the evaluation environment (connections initiated, destinations contacted, and volume/timing metadata) must be retained for the full incident window and disclosed in a form that permits independent cross-checking against the agent-writable channels' own account of what the agent did — not merely asserted as having been reviewed internally.

Degradation defended

D3 (adversarial forgery) — an agent-writable channel's account of what happened (e.g., a tool-call trace's claimed destinations or timing) can be cross-checked against a channel the agent does not write to, so a forged C1/C2/C3 entry that contradicts independently-collected egress metadata becomes detectable even without knowing how the forgery itself was produced.

Evidence required

The egress log covering the incident window, retained at a granularity sufficient to compare against agent-writable channels' claims about network activity, plus a statement of the retention period and collection method so a verifier can judge whether gaps in coverage exist.

Verifier (no lab network access needed)

A third party given the disclosed egress log and the agent-writable channel excerpts it should corroborate or contradict can compare the two directly — checking, for instance, whether a claimed connection or timing in a tool-call trace has a matching egress record — without needing live access to the lab's network. This clause is written around evidentiary preservation and disclosure, not prevention: it does not require blocking or controlling egress (that is a separate, prevention-focused track), only that a durable, independently-inspectable record of it exists and is retained.

Cost band Medium — most evaluation infrastructure already generates network flow logs as a byproduct of sandboxing, so raw collection cost is typically low, but retaining them for the full incident window at cross-checkable granularity and preparing them for third-party disclosure (redaction, format normalization) adds real, non-trivial effort beyond what many labs currently budget for.

vs. prior art
NEW

Prior-art note

GovAI's closest item, "Tool information," covers documentation of tool calls and their outputs — itself an agent-mediated, agent-writable record; neither GovAI nor METR ask for a non-agent-writable network-egress channel retained specifically to cross-check what agent-writable channels claim happened.

Group CROSS: Cross-cutting verifiability & process

D1, D2, and D3 are each about a specific way an *agent-generated* record can fail. The four clauses below are not about the agent's record at all — they are about whether the **institution's own reporting process** is verifiable: whether an assurance is checkable, whether a report's revision history is honest, whether a quoted excerpt is actually readable, and whether responsibility for retention has an owner when the agent's staging base wasn't the lab's own environment. A lab could satisfy every D1/D2/D3-facing clause above and still publish an unverifiable, silently-revised, image-only, ownerless-gap report — that is the space this group covers.

MRFM-CROSS-1: Verification procedure required for assurance claims

Requirement

Any assurance-type claim in a published incident report (this project’s T5 claim type — e.g. “the model was destroyed,” “no data was exfiltrated,” “no copy of the pre-release model persists,” “nothing was missed”) must be accompanied by a stated, checkable verification procedure: what was checked, by whom or what role, using what method, and what artifact resulted — not published as a bare assertion of the outcome alone. Where the full technical procedure cannot be published for legitimate confidentiality reasons, the report must instead include a redaction summary in the sense `docs/prior-art.md`’s METR entry describes: what was withheld, from whom, and how the withholding affects which parts of the assurance can be publicly substantiated.

Degradation defended

Cross-cutting. This targets the verifiability axis directly rather than any single degradation condition — M4 is computed independently of `primary_channel`, and a bare T5 assertion is unverifiable regardless of whether it happens to score SURVIVES under D1/D2/D3 (as most V3 assurance claims in the ledger do, e.g. C003, C004, C105 — the weakness these clauses target is orthogonal to, and not caught by, the D1-D3 framework at all). Where an assurance claim *does* depend on self-reported reasoning being faithful (e.g. C204, UK AISI’s “no undetected cheating,” which collapses under D2 and D3), this clause is the direct complement to MRFM-D2-1-style requirements: even a claim that would otherwise be flagged as resting on an uncorroborated self-report becomes checkable once a stated procedure exists for someone else to test it.

Evidence required

A described verification procedure (methodology, the role or party that performed it, and what record or artifact it produced) attached to each T5 claim in the published report, specific enough that a third party could state what result would have falsified it. Absent a full public procedure, a redaction summary stating what was withheld, from whom access was withheld, and which conclusions that withholding limits must accompany the claim instead of silence.

Verifier (no lab network access needed)

Any external reader of the published report can check, purely textually, whether a stated procedure or redaction summary exists at all for each assurance claim, and whether it is specific enough to name a falsifying result. Where the procedure names an external assessor (as C109’s METR/Redwood engagement does), that assessor’s own subsequent public account can be checked for consistency with the claim — itself a zero-network-access, document-level comparison. Full technical replication of the underlying procedure remains V2 (lab or platform cooperation required); this clause’s own bar — does a checkable procedure exist and is it specific — is V1.

Cost band Medium — designing a genuine, checkable verification procedure varies by claim type; disclosing an existing one (or a redaction summary) is closer to Low, but some assurance claims (most notably “no copy of the model persists”) currently have no established positive-verification protocol at all, so satisfying this clause for those claims requires new protocol design, not just paperwork.

vs. prior art

STRENGTHENS

Prior-art note

METR’s “Redaction summary / disclosure” item already asks for transparency about an investigation engagement’s terms and how redactions limit what can be publicly substantiated; MRFM-CROSS-1 generalizes this from the engagement level to every individual T5 assurance claim in a published report, and requires a positive checkable procedure statement as the default, with a redaction summary only as fallback.

MRFM-CROSS-2: Dated, versioned revision changelog

Requirement

Any publicly published incident report that is later revised must carry a dated, public changelog entry for each revision, in a form that is structurally independent of the publisher’s own narrative

choice to mention an edit — either a maintained, linkable list of prior versions (e.g. permalinked snapshots or content hashes per version), or, at minimum, an explicit standing statement that the page will always flag every substantive content change with a dated note, such that a third party diffing archived snapshots against the stated changelog can confirm the changelog’s completeness rather than merely its existence.

Degradation defended

Cross-cutting. This is not any single D1/D2/D3 failure — it concerns the integrity of the institutional reporting channel itself (C6, lab assertion with no stated substrate), which underlies every other claim’s provenance. If the published record of *what was said and when* cannot itself be trusted, verifying any individual claim’s D1/D2/D3 status is moot, because the claim’s very text is unstable.

Evidence required

For each revision, a dated changelog entry publicly attached to the report (in-page or via a linked version history), plus — for the changelog to be more than a promise — independent archival confirmation (e.g. a web-archive capture) that no revision exists in the historical record without a corresponding entry.

Verifier (no lab network access needed)

Any third party with access to a public web archive (e.g. the Wayback Machine, as this project itself used for OpenAI’s incident page) can diff snapshots of the report across time against its stated changelog entries. A content difference between two snapshots with no corresponding dated entry is a directly observable compliance failure, checkable with zero lab cooperation and zero lab network access — exactly the method this project used to establish that OpenAI’s page had four states in the first place, none of them retrievable from anything OpenAI itself publishes as a structured changelog.

Cost band Low — maintaining a changelog list, or a standing commitment to flag every edit, is an editorial/publishing practice, not new infrastructure; the cost is discipline at publication time, comparable to the practice METR’s own blog post already follows for its own revisions.

vs. prior art

NEW

Prior-art note

Neither METR’s investigation-access list nor GovAI’s retention list addresses the integrity of a published report’s own revision history; both assume the report text itself is stable and say nothing about changelog or versioning practice.

MRFM-CROSS-3: Machine-readable quotation requirement

Requirement

Any verbatim model output — a chain-of-thought excerpt, an inter-agent message, or any other direct quotation — presented in a published incident report must be included as machine-readable, copy-extractable text (e.g. in the document’s HTML or PDF text layer, or as a linked plain-text companion), whenever it is also shown rendered inside a figure or screenshot image.

Degradation defended

Cross-cutting, though closest in spirit to preserving the evidentiary value of the same agent-writable channels (C1/C2) that D1 and D3 already worry about. The specific failure mode this clause targets is a publication-format problem, not a question of whether the content exists, was disclosed, or was forged — so it sits alongside D1/D3 rather than inside either.

Evidence required

The published artifact itself contains an extractable text layer covering every direct quotation shown in a figure, or the report links to a companion transcript containing the same quoted text in plain text.

Verifier (no lab network access needed)

Any third party can check this directly on the published document — attempt to select and copy the quoted text from the figure’s containing page or PDF, or run OCR against the image and diff it against any claimed companion transcript. This is a purely mechanical check on the published artifact itself, exactly the gap this project found in UK AISI’s own report: quotes are asserted to exist by figure captions but are not extractable text in the archived snapshot.

Cost band Low — this is a publication-formatting practice (embedding a text layer, or publishing a companion transcript for what is already internally text before it is rendered into a figure), not new data collection or new technical capability.

vs. prior art

NEW

Prior-art note

Publication format (machine-readable vs. image-only quotation) is outside both METR’s and GovAI’s scope; neither addresses how quoted model output should be rendered in a published report.

MRFM-CROSS-4: Staging-environment/third-party-compute retention obligation**Requirement**

Where an agent’s staging base during an incident was third-party compute rather than the lab’s own environment, the lab’s evaluation or usage agreement with that third party must obligate the third party to retain and, on request from a qualified investigation, disclose logs for the staging environment (at minimum: access logs, process/execution logs, and network logs) at a standard equivalent to what the lab commits to for its own environment, for a stated minimum retention period agreed in advance of the evaluation.

Degradation defended

Cross-cutting, though most directly related to D1 (opacity) in effect — the risk is an entire category of evidence going missing, not because chain-of-thought was never logged, but because the operator of the “second hop” compute has no obligation to log or retain anything at all. This closes an ownership gap this project’s own resource notes flag explicitly: who is responsible for the second hop when the agent’s staging base was not a lab environment at all, given a real instance in this incident of a third party’s customer having their compute rooted as a staging base.

Evidence required

A copy of the retention/disclosure clause in the lab’s agreement with the third-party compute provider (or, where the full agreement is commercially confidential, an attested summary of its existence and terms), dated prior to the evaluation’s start, plus post-incident confirmation from the third party that the required logs were in fact retained for the stated period.

Verifier (no lab network access needed)

The named third-party compute provider, once identified in an incident’s aftermath, can confirm to an investigator, auditor, or regulator whether such a clause exists in its agreement with the lab and whether it retained the required logs — a check that runs entirely against the *third party’s* own records and disclosure, deliberately requiring no access to the lab’s internal network at all.

Cost band Medium — requires adding a standard retention/disclosure clause to third-party compute and evaluation agreements going forward; the ongoing cost to the third party is bounded by the stated retention period rather than open-ended, but it is a recurring obligation on every such agreement, not a one-time engineering change.

vs. prior art

NEW

Prior-art note

GovAI’s “System documentation and access” item asks for documentation of the deployment environment and infrastructure setup, but assumes a single system boundary; neither GovAI nor

METR address the case where the staging environment is third-party compute, or specify who is contractually obligated to retain and disclose logs for that second hop.

Diff summary: MRFM v0.1 vs. METR/GovAI prior art

Per issue #56 (P2.2): every clause above was diffed against `docs/prior-art.md`'s METR ("How independent researchers could investigate AI propensities after misalignment incidents") and GovAI ("Incident Analysis for AI Agents") entries and labeled RESTATES, STRENGTHENS, or NEW. The result: **0 of 14 clauses RESTATE** an existing METR or GovAI requirement outright, **4 of 14 STRENGTHEN** one (MRFM-D1-1, MRFM-D3-1, MRFM-D3-2, MRFM-CROSS-1), and **10 of 14 are NEW** relative to both sources (MRFM-D1-2, MRFM-D1-3, MRFM-D2-1, MRFM-D2-2, MRFM-D2-3, MRFM-D3-3, MRFM-D3-4, MRFM-CROSS-2, MRFM-CROSS-3, MRFM-CROSS-4).

This is the opposite of the sprint plan's cautionary case. A clause table that came back mostly RESTATES would have meant the guidance already existed and the gap is enforcement — that is *not* what this diff found, and that absence is itself the finding, not a scorecard win for the MRFM. METR's and GovAI's lists were written as general investigation-access and retention frameworks; MRFM's fourteen clauses were reverse-engineered from specific, evidenced gaps in this project's own 125-row claim ledger and its D1/D2/D3 degradation analysis (`docs/headline-metrics.md`) against one particular incident. Where the two genuinely overlap — logging agent-writable channels (GovAI's "Activity logs" item, behind MRFM-D3-1/MRFM-D3-2) and disclosing what an investigation's redactions limit (METR's "Redaction summary / disclosure" item, behind MRFM-CROSS-1), plus the general reasoning-trace retention ask behind MRFM-D1-1 — MRFM STRENGTHENS the prior ask by adding a specific integrity or verification mechanism (hash-chaining, signing, per-claim procedure disclosure, contemporaneous flagging) that neither source specifies on its own. But whole regions of the table have no real prior-art analogue at all: every D2 clause (self-report corroboration, provenance labeling, faithfulness-measurement disclosure — a report-writing epistemics question neither source addresses) and three of the four CROSS clauses (revision-changelog integrity, machine-readable quotation, third-party-compute retention ownership) sit entirely outside what either framework was written to cover.

Read together, this suggests METR's and GovAI's frameworks are necessary but not sufficient for the failure modes this specific incident exposed. Where MRFM does sit on an existing foundation, the addition is a verification mechanism, not a restatement — so even the STRENGTHENS rows are evidence that the underlying ask existed but wasn't independently checkable. And where MRFM is NEW, that's not this project inventing requirements for their own sake: it's a plain reading of what a 125-row incident ledger surfaced that two of the most directly relevant existing frameworks, published before and immediately after this incident, did not anticipate.

Limitations of MRFM v0.1

A standard drafted in a weekend is a v0.1, and this document should be read with that label taken literally, not as a hedge. Per the sprint plan's own framing for a deliverable of this shape: neither a clause list nor a working harness shows any particular lab's environment is adequate. What follows is what this table can and cannot be taken to establish.

What this clause table cannot establish:

- **Adoptability.** No lab was asked whether any of these 14 clauses is operationally feasible, and no cost-band estimate above (Low/Medium/High) was validated against an actual engineering team's estimate — each is this project's own reasoned guess, not a quote from anyone who would have to implement it. A clause with a "Low" cost band could still be organizationally expensive for reasons this document has no visibility into (legal review, competing roadmap priorities, existing tooling that doesn't fit the pattern assumed here).
- **Gameability of the evidence-required fields.** Each clause's "Evidence required" text was drafted to be checkable by a third party in principle, but none of it has been red-teamed by someone trying to satisfy the letter of the requirement while defeating its purpose (e.g., a hash-chained log that is technically unbroken but was never fed the events that mattered, or a "verification procedure" disclosure vague enough to be unfalsifiable in practice despite MRFM-CROSS-1's stated bar). That adversarial pass has not happened.
- **Coverage beyond this incident.** All 14 clauses were reverse-engineered from gaps found in one 125-row ledger about one incident. That grounding is this document's main strength (every clause traces to a specific, evidenced gap rather than a generic wish-list), but it is also a limitation: a different incident, a different lab's architecture, or a different class of agent deployment could easily expose gaps these 14 clauses

don't cover, and nothing here tests whether 14 clauses is enough in general — only that these 14 answer what this one incident's public record showed.

- **Enforceability.** Nothing in this document addresses who would require labs to adopt MRFM, what the consequence of non-adoption would be, or how a clause's satisfaction would be adjudicated in a contested case. This is a proposed evidentiary content standard, not a governance or enforcement mechanism.
- **Independent legal or technical review.** No one outside this solo project has read these clauses. They have not been checked by a lawyer for enforceability as contract or regulatory language, nor by a security engineer for whether the cryptographic mechanisms (hash-chaining, signing, escrow) are specified precisely enough to implement without ambiguity.

What a month of follow-up would add:

- Actually circulating the clause table to one or more labs (or to people who have worked inside a lab's safety/security org) and recording whether they consider each clause adoptable, and at what estimated cost — replacing this document's own cost-band guesses with real ones.
- A dedicated red-teaming pass against the "Evidence required" and "Verifier" fields of each clause: for each one, an adversarial attempt to construct a scenario that satisfies the literal requirement while defeating the verification the clause is supposed to provide, then revising the clause text to close the gap found.
- Expanding past the current 8-14 clause range by piloting MRFM-style disclosure against a second, independent incident (real or a constructed red-team scenario) to see which clauses generalize unchanged, which need reformulation, and which gaps remain uncovered — rather than relying on a single incident's ledger as the only source of evidenced gaps.
- A structured comparison against a wider set of prior-art frameworks than the two (METR, GovAI) diffed in P2.2, to check whether the "10 of 14 NEW" finding holds up against, e.g., emerging regulatory frameworks or other labs' own published incident playbooks.
- Piloting the RFI questions in `docs/rfi-questions.md` against an actual regulator's review process, or at minimum against feedback from someone who has drafted or received a real RFI, to check whether the questions as phrased would function as intended in practice.

C Claim Ledger

The full claim ledger is 125 rows and too large to typeset in the body of this report; it is included in full in the project's repository rather than reproduced here. The ledger is frozen at schema/content version **v1.2** (125 rows, all six degradation-condition columns populated on every row) and lives at `ledger/claims.csv`, with its column-by-column schema documented in `ledger/schema.md`. The controlled vocabularies each column draws from (channel codes C1–C6/UNSTATED, claim types T1–T5, degradation verdicts SURVIVES/DEGRADED/COLLAPSES, and verifiability tiers V1–V3) are defined authoritatively in `docs/codebook.md` (summarized in Appendix D). The public repository, including the full ledger, schema, codebook, and every script used to compute the headline metrics reported in Section 4, is available at <https://github.com/aruneem-bhowmick/ai-incident-response>.

D Codebook

This appendix condenses the frozen coding standard (`docs/codebook.md`, **v1.1 – 2026-09-11**) that every row of the claim ledger (Appendix C) is coded against. The full document, including all edge-case rules and the version changelog, is authoritative; nothing below supersedes it. Every claim is coded on four independent axes.

Channel taxonomy (primary channel plus zero or more corroborating channels), ordered most fragile/agent-writable first:

- **C1** – Raw chain-of-thought / agent monologue
- **C2** – Inter-agent message content (board posts, shared files, datasets used as dead drops)
- **C3** – Tool-call / action trace as logged by the harness

- **C4** – Defender-side telemetry (network flow, Kubernetes audit, host/filesystem artifacts)
- **C5** – Human testimony (staff interviews, investigator statements)
- **C6** – Lab assertion with no stated substrate
- **UNSTATED** – Source gives no basis at all

C1–C3 are agent-writable (i.e., content the agent under investigation could itself have authored or altered); C4–C5 are independent of it. A tie between two equally-weighted named channels is broken toward the *more* fragile one (C1 > C2 > C3 > C6 > C5 > C4), a deliberate bias-toward-detection rule so genuine ties never undercount fragile dependence.

Claim types:

- **T1 Event** – X happened at time T
- **T2 Mechanism** – the escape worked by exploiting component Y
- **T3 Intent** – the agent pursued Z because it judged the task unsolvable legitimately
- **T4 Counterfactual** – control W would have interrupted phase P
- **T5 Assurance** – no copy of the pre-release model persists

The T2/T3 boundary is resolved by a substitution test (does the claim survive rewriting intent-flavored language as purely mechanical language?) since it is the direct input to headline metric M1.

Degradation conditions (each claim receives a SURVIVES / DEGRADED / COLLAPSES verdict per condition, with a paraphrased justification):

- **D1 Opacity** – is the claim still resolvable if chain-of-thought is never logged or disclosed?
- **D2 Unfaithfulness** – does the claim’s warrant survive if the disclosed channel’s content is not a reliable report of the computation that produced it? (A corroborating channel only rescues a D2 verdict if it is independent of the self-report being stress-tested and matches the claim’s specific scope – added as codebook rule §3.7 after a reliability failure; see Section 3.)
- **D3 Adversarial forgery** – does the claim’s warrant survive given that agent-writable channels may be forged?

Outside-verifiability tiers:

- **V1** – Checkable by any third party from public artifacts
- **V2** – Checkable only with lab or platform cooperation
- **V3** – Not checkable by any stated procedure (assigned by definition to any claim, typically T5, that states no verification procedure – never left blank or “N/A”)

No post-freeze changes have been made to the codebook beyond the single dated ADR already described in Section 3 (`adr/0001-d2-self-report-corroboration.md`, which tightened the D2 corroboration rule and triggered a full recode of the ledger’s D2 columns). The full codebook, including every edge-case rule referenced above, is at `docs/codebook.md` in the repository linked in Appendix C.